

# The Art of Iterating: Update-Strategies in Agent-Based Simulations

Jonathan THALER

January 24, 2017

## Abstract

When developing a model for an Agent-Based Simulation (ABS) it is of very importance to select the right update-strategy for the agents to produce the desired results. In this paper we develop a systematic treatment of all general update-strategies in ABS and discuss their philosophical interpretation and meaning. Further we discuss the suitability of the very three different programming languages Java, Scala with Actors and Haskell to implement each of the update-strategies. Thus this papers contribution is the development of a general terminology of update-strategies and their implementation issues in various kinds of programming languages.

## 1 Introduction

In the paper of [5] the authors showed that the results of the simulation of the classic prisoners-dilemma on a 2D-grid reported in [6] depends on a a very specific strategy of iterating this simulation and show that the beautiful patterns as reported by [6] will not form when selecting a different iteration-strategy. Although the authors differentiated between two strategies, their description still lacks precision and generality which we will try to repair in this paper. Although they too discussed philosophical aspects of choosing one strategy over the other, they lacked to generalize their observation. We will do so in the central message of our paper by stressing that when doing Agent-Based Simulation & Modelling (ABM/S) *it is of most importance to select the right iteration-strategy which reflects and supports the corresponding semantics of the model*. We find that this awareness is yet still under-represented in the literature of ABM/S and most important of all is lacking a systematic treatment. Thus our contribution in this paper is to provide such a systematic treatment by

- Presenting all the general iteration-strategies which are possible in an ABM/S.
- Developing a systematic terminology of talking about them.

- Giving the philosophical interpretation and meaning of each of them.
- Comparing the 3 programming languages Java, Haskell and Scala in regard of their suitability to implement each of these strategies.

Besides the systematic treatment of all the general iteration-strategies the paper presents another novelty which is its inclusion of the pure functional declarative language Haskell in the comparison. This language has so far been neglected by the ABM/S community which is dominated by object-oriented (OO) programming languages like Java thus the usage of Haskell presents a real, original novelty in this paper.

## 2 Related Research

TODO: [5] where the first to discuss the differences update-strategies can make  
 TODO: [1] give an approach for ABM/S on GPUs. a seemingly completely different form but which, we hypothesize, can be roughly mapped to our PAR-Strategy.

[3] sketch a minimal agent-framework in Haskell which is very similar in the basic structure of ours, also utilizing an agent-transforming function which consumes incoming messages and produces outgoing ones. This proves that this approach, very well developed in ABM/S, seems to be a very natural one also to apply to Haskell. Their focus is more on economic simulations and instead of iterating a simulation with a global time, their focus is on how to synchronize agents which have internal, local transition times. They introduce a time-keeper agent which synchronizes the actions of all of the agents thus we argue that our framework is able to capture it faithfully using the *Actor-Strategy* utilizing either a timer-keeper as they do or through the access of the global shared-environment. Although their work uses Haskell as well, this does not diminish the novelty of our approach using Haskell because our focus is very different from them.

TODO: [8]: Master-Thesis "The Agent-based Simulation Environment in Java" by Yuxuan Jin under supervision of Peer-Olaf

TODO: [4]: comparison of c++ and java for ABM/S

## 3 The Problem

ABM/S is a method of modelling and simulating a system where the global behaviour may be unknown but the behaviour and interactions of the parts making up the system is of knowledge [7]. Those parts, called Agents, are modelled and simulated out of which then the aggregate global behaviour of the whole system emerges. Thus the central aspect of ABM/S is the concept of an Agent which can be understood as a metaphor for a unique pro-active unit, able to spawn new Agents, interacting with other Agents in a network of neighbours by exchange of messages which are situated in a generic environment. Thus we informally assume the following about our Agents:

- They have a unique identifier
- They can initiate actions on their own e.g. change their own state, send messages, create new agents, kill themselves,...
- They can react to messages they receive with actions (see above)
- They can interact with a generic environment they are situated in

An implementation of an ABS must solve thus solve three fundamental problems:

1. Source of pro-activity  
How can an Agent initiate actions without external stimuli?
2. Message-Processing  
When is a message  $m$ , visible to Agent  $B$ , processed by it?
3. Semantics of Message-Delivery  
When is a message  $m$ , sent by Agent  $A$  to Agent  $B$ , visible to  $B$ ?

In computer systems, pro-activity, the ability to initiate actions on its own without external stimuli, is only possible when there is some internal stimuli, most naturally represented by some generic notion of monotonic increasing time-flow. This can either be some physical real-time system-clock which counts the milliseconds since 1970 (thus binding the time-flow of the system to the one of the 'real-world') or a virtual simulation-clock which is just a monotonic increasing natural number. As we are in a discrete computer-system, this time-flow must be discretized as well in discrete steps and each step must be made available to the Agent, acting as one internal stimuli. This allows the Agent then to perceive time and become pro-active depending on time (NOTE: we could argue that this is not really pro-activity because it depends always on time, but there is really now other way of doing this in our current implementation of computer-systems.). Independent of the representation of the time-flow we have the two fundamental choices whether the time-flow is local to the Agent or whether it is a system-global time-flow.

The semantics of message-delivery define when sent messages are visible to the receivers so they can process them and react to them. The only two ways of implementing them are that messages are visible either *immediately* or after a synchronization point between the sender and receiver. Such a synchronization point can be a local one, just between the two or a global one between all Agents in the system.

Basically we can say that we want to process a message as soon as it is visible to us but this is not how real computer-systems can work. In a real system each Agent would have a message-box into which the messages are posted so the Agent can then check its mail-box for new messages. The question is then when the Agent is going to poll for new messages? Clearly what we need is

a recurring, regular trigger which allows the Agent to poll the mail-box and process all queued messages. We argue that the most natural approach is to bind this trigger to the time-flow step which provides pro-activity.

To solve these problems an update-strategy is implemented which will iterate through the Agents in *some* way and allow the Agents to perform these steps. It is immediately clear that different choices in the specific problems will lead to different system behaviour. To discuss this we will first present all possible update-strategies and their details in the next section and then outline how they influence the system behaviour.

## 4 Update-Strategies

2 Pages

In this section we will present all the update-strategies which are available in ABM/S in a general form, discuss abstract details independent from programming languages, give philosophical meaning and interpretation of them and advices for selecting them for which model.

<i>Name</i>	<i>Time-Flow</i>	<i>Iteration-Order</i>	<i>Deterministic</i>
<b>Seq</b>	Global	Sequential	Yes
<b>Par</b>	Global	Parallel	Yes
<b>Con</b>	Global	Concurrent	No
<b>Act</b>	Local	Actor	No

Table 1: List of all general update-strategies in ABM/S.

### 4.1 Seq

**Description:** This strategy has a global time-flow and in each time-step iterates through all the agents and updates one Agent after another. Messages sent and changes to the environment made by Agents are visible immediately. More formally, we assume that, given the updates are done in order of the index  $i = [1..n]$ , then Agents  $a_{n>i}$  see the changes to environment and messages sent to them by Agent  $a_i$ . Note that there is no source of randomness and non-determinism thus rendering this strategy to be completely deterministic in each step.

**Time-Updates:** The most straight-forward approach is to keep the time constant for each agent in one iteration. This though may seem non-logic as the actions of preceding Agents are visible to later ones but time has not changed since then. Thus if the model semantics require that time changes with every Agent we could advance for every agent by fraction of  $dt$ :  $\text{agent-time} = t + (a_i * dt/n)$  where  $t$  is the current simulation time,  $a_i$  the agents index,  $dt$  the amount of time the simulation will be advanced by and  $n$  the number of agents. In the end the new current time will be then  $t_{\text{next}} = t_{\text{curr}} + dt$ . Other possibilities

of advancing is  $\text{agent-time} = t + a_i * dt$ . in the end the new current time will be then  $t_{\text{next}} = t_{\text{curr}} + n * dt$

**Extensions:** If the sequential iteration from 1..n imposes an advantage over the Agents further ahead or behind in the queue (e.g. if it is of benefit when making choices earlier than others in auctions or later when more information is available) then one could use random-walk iteration where in each time-step the agents are shuffled before iterated. Note that although this would introduce randomness in the model the source is a random-number generator thus reproduce-able.

## 4.2 Par

**Description:** This strategy has a global time-flow and in each time-step iterates through all the agents and updates all Agents in parallel. Messages sent and changes to the environment made by Agents are visible in the next global step. We can think about this strategy that all Agents make their moves at the same time.

**Environment:** If one wants to change the environment in a way that it would be visible to other Agents this is regarded as a systematic error in this strategy. First it is not logical because all actions are meant to happen at the same time and also it would implicitly induce an ordering thus violating the *happens at the same time* idea. Thus we require different semantics for accessing the environment in this strategy. We introduce thus a *global* environment which is made up of the set of *local* environments. Each local environment is owned by an Agent thus there are as many local environments as there are Agents. The semantics are then as follows: in each step all Agents can *read* the global environment and *read/write* their local environment. The changes to a local environment are only visible *after* the local step and can be fed back into the global environment after the parallel processing of the Agents.

**Time-Updates:** Time really stays constant in this case for all Agents in one step as all updates happen really at the same *virtual* time. It would make no sense to advance the time between Agents as all actions are meant to happen at the same time, without imposing any ordering amongst them.

**Semantics:** It does not make a difference if the Agents are really computed in parallel or just sequentially, due to the semantics of changes, this has the same effect. In this case it will make no difference how we iterate over the agents (sequentially, randomly), the outcome *has to be* the same - it is event-ordering invariant as all events/updates happen *virtually* at the *same time*. Thus if one needs to have the semantics of writes on the whole (global) environment in ones model, then this strategy is not the right one and one should resort to one of the other strategies.

## 4.3 Con

**Description:** This strategy has a global time-flow and in each time-step iterates through all the agents and updates all Agents in parallel but all messages sent

and changes to the environment are immediately visible. Thus this strategy can be understood as a mix of Seq and Par: all Agents run at the same time with actions becoming immediately visible.

**Semantics:** It is important to realize that, when running Agents in parallel which are able to see actions by others immediately, this is the very definition of concurrency: parallel execution with mutual read/write access to shared data. Of course this shared data-access needs to be synchronized which in turn will introduce event-orderings in the execution of the Agents. Thus at this point we have a source of inherent non-determinism: although we would ignore any hardware-model of concurrency at some point we need arbitration to decide which Agent gets access first to a shared resource thus arriving at non-deterministic solutions - this will become much clearer in the results-section. This has the very important influence that repeated runs with the same configuration of the Agents and the Model may lead to different results each time.

#### 4.4 Actors - Strategy

TODO: discuss how local-time can be handled: real-time or simulation-time - its always local and not synchronized globally because then we would end up in Concurrent Strategy

in the Act-version we need to observe the agents: we need to sample them regularly =, we have all the issues with sampling

In this case there is no global iteration over steps but all the Agents run in parallel, doing local stepping and communicate with each other either through shared state or messages. Note that this does not impose any specific ordering of the update and can thus regarded to be real random due to its concurrent nature. It is possible to simulate the global-stepping methods from above by introducing some global locking forcing the agents into lock-step. This is the approach chosen for Scala & Actors.

**Semantics:** This is the most general one of all the strategies as it can emulate all the others by introducing the necessary synchronization mechanisms and Agents.

#### 4.5 Update-Strategies

1. All states are copied/frozen which has the effect that all agents update their positions *simultaneously*
2. Updating one agent after another utilizing aliasing (sharing of references) to allow agents updated *after* agents before to see the agents updated before them. Here we have also two strategies: deterministic- and random-traversal.
3. Local observations: Akka

## 4.6 Different results with different Update-Strategies?

Problem: the following properties have to be the same to reproduce the same results in different implementations:

Same initial data: Random-Number-Generators Same numerical-computation: floating-point arithmetic Same ordering of events: update-strategy, traversal, parallelism, concurrency

- Same Random-Number Generator (RNG) algorithm which must produce the same sequence given the same initial seed.
- Same Floating-Point arithmetic
- Same ordering of events: in Scala & Actors this is impossible to achieve because actors run in parallel thus relying on os-specific non-deterministic scheduling. Note that although the scheduling algorithm is of course deterministic in all os (i guess) the time when a thread is scheduled depends on the current state of the system which can change all the time due to *very* high number of variables outside of influence (some of the non-deterministic): user-input, network-input, .... which in effect make the system appear as non-deterministic due to highly complex dependencies and feedback.
- Same dt sequence =  $\Delta t$  dt MUST NOT come from GUI/rendering-loop because gui/rendering is, as all parallelism/concurrency subject to performance variations depending on scheduling and load of OS.

It is possible to compare the influences of update-strategies in the Java implementation by running two exact simulations (agentcount, speed, dt, herodistribution, random-seed, world-type) in lock-step and comparing the positions of the agent-pairs with same ids after each iteration. If either the x or y coordinate is not equal then the positions are defined to be *not* equal and thus we assume the simulations have then diverged from each other.

It is clear that we cannot compare two floating-point numbers by trivial `==` operator as floating-point numbers always suffer rounding errors thus introducing imprecision. What may seem to be a straight-forward solution would be to introduce some epsilon, measuring the absolute error:  $\text{abs}(x1 - x2) < \text{epsilon}$ , but this still has its pitfalls. The problem with this is that, when number being compared are very small as well then epsilon could be far too big thus returning to be true despite the small numbers are compared to each other quite different. Also if the numbers are very large the epsilon could end up being smaller than the smallest rounding error, so that this comparison will always return false. The solution would be to look at the *relative error*:  $\text{abs}((a-b)/b) < \text{epsilon}$ .

The problem of introducing a relative error is that in our case although the relative error can be very small the comparison could be determined to be different but looking in fact exactly the same without being able to be distinguished with the eye. Thus we make use of the fact that our coordinates are virtual

ones, always being in the range of  $[0..1]$  and are falling back to the measure of absolute error with an epsilon of 0.1. Why this big epsilon? Because this will then definitely show us that the simulation is *different*.

The question is then which update-strategies lead to diverging results. The hypothesis is that when doing simultaneous updates it should make no difference when doing random-traversal or deterministic traversal =; when comparing two simulations with simultaneous updates and all the same except first random- and the other deterministic traversal then they should never diverge. Why? Because in the simultaneous updates there is no ordering introduced, all states are frozen and thus the ordering of the updates should have no influence, *both simulations should never diverge, independent how  $dt$  and epsilon are selected*.

Do the simulation-results support the hypothesis? Yes they support the hypothesis - even in the worst cast with very large  $dt$  compared to epsilon (e.g.  $dt = 1.0$ , epsilon = 1.0-12)

The 2nd hypothesis is then of course that when doing consecutive updates the simulations will *always* diverge independent when having different traversal-strategies.

Simulations show that the selection of  $dt$  is crucial in how fast the simulations diverge when using different traversal-strategies. The observation is that *The larger  $dt$  the faster they diverge and the more substantial and earlier the divergence..* Of course it is not possible to proof using simulations alone that they will always diverge when having different traversal-strategies. Maybe looking at the dynamics of the error (the maximum of the difference of the x and y pairs) would reveal some insight?

The 3rd hypothesis is that the number of agents should also lead to increased speed of divergence when having different traversal-strategies. This could be shown when going from 60 agents with a  $dt$  of 0.01 which never exceeded a global error of 0.02 to 6000 agents which after 3239 steps exceeded the absolute error of 0.1.

## 4.7 Reproducing Results in different Implementations

actors: time is always local and thus information as well. if we fall back to a global time like system time we would also fall back to real-time. anyway in distributed systems clock sync is a very non-trivial problem and inherently not possible (really?). thus using some global clock on a metalevel above/outside the simulation will only buy us more problems than it would solve us. real-time does not help either as it is never hard real time and thus also unpredictable: if one tells the actor to send itself a message after 100ms then one relies on the capability of the OS-timer and scheduler to schedule exactly after 100ms: something which is always possible thus 100ms are never hard 100ms but soft with variations.



qualitative comparison: print picture with patterns. all implementations are able to reproduce these patterns independent from the update strategy no need to compare individual runs and waste time in implementing RNGs, what is more interesting is whether the qualitative results are the same: does the system show the same emergent behaviour? Of course if we can show that the system will behave exactly the same then it will also exhibit the same emergent behaviour but that is not possible under some circumstances e.g. the simulation-runs of Akka are always unique and never comparable due to random event-ordering produced by concurrency & scheduling. Also we don't have to proof the obvious: given the same algorithm, the same random-data, the same treatment of numbers and the same ordering of events, the outcome *must* be the same, otherwise there are bugs in the program. Thus when comparing results given all the above mentioned properties are the same one in effect tests only if the programs contain no bugs - or the same bugs, if they *are the same*.

Thus we can say: the systems behave qualitatively the same under different event-orderings.

Thus the essence of this boils down to the question: "Is the emergent behaviour of the system is stable under random/different/varying event-ordering?". In this case it seems to be so as proofed by the Akka implementation. In fact this is a very desirable property of a system showing emergent behaviour but we need to get much more precise here: what is an event? what is an emergent behaviour of a system? what is random-ordering of events? (Note: obviously we are speaking about repeated runs of a system where the initial conditions may be the same but due to implementation details like concurrency we get a different event-ordering in each simulation-run, thus the event-orderings vary between runs, they can be in fact be regarded as random).

## 5 Results

5 Pages

This is now very programming-language specific

- Mapping the strategies to 3 programming-languages: Java, Scala with Actors, Haskell
- Comparing the programming languages in regard of their suitability to implement each of these strategies
- Screen-shots of results of the same simulation-model with all the strategies

### 5.1 Selection of the Languages

-j Java: supports global data =j suitable to implement global decisions: implementing global-time, sequential iteration with global decisions -j Haskell: has no global data =j local decisions (has support for global data through STM/IO

but then loses very power?) =, implementing global-time, parallel iteration with local-decisions. -, Haskell STM solution =, implementing concurrent version using STM? but this is very complicated in its own right but utilizing STM it will be much more easier than in java -, Scala: mixed, can do both =, implementing local time with random iteration and local decisions

Note that we didn't select a language where the memory-management falls in the hands of the developer. This would be the case e.g. in C++. This was looked into partially by [4] but the focus of this paper is not on this issue as it would complicate things dramatically. All used languages are garbage-collected / the developer does not need to care how memory is cleared up.

## 5.2 Selected Models

### 5.2.1 Heroes & Cowards

### 5.2.2 SIRS

### 5.2.3 Wildfire

### 5.2.4 Spatial Game

## 5.3 Java

sequential more natural in java, parallel needs to "think functional" in java concurrency and actors always difficult in java despite java provides very good synchronization primitives

IMPORTANT: no need to explicitly add the environment, as this can be individually implemented by the agents on a shared basis (references)

sequential is able to work completely without messages and only by accessing references to neighbours but we explicitly don't want to follow this obvious way and stick to the send/receive message paradigm. we keep Agent-instance references

parallel then needs to utilize messages because would violate the parallel implementation. TODO: split state from agent and only update agent-state problem: can send references through messages: share data although the interfaces encourage it we cannot prevent the agents to use agent-references and directly accessing. a workaround would be to create new agent-instances in every iteration-step which would make old references useless but this doesn't protect us from concurrency issues with a current iteration (the copying must take place within a synchronized block, thus implicitly assuming ordering, something we don't want) and besides, can always work around and update the references. that's the toll of side-effects: faster execution but less control over abuse tried to clone agents in each step and let them collect their messages =, extremely slow  
conc: expect it to be a pain in the ass with java. it is not: it's the same interface as in SEQ with updates running parallel like in PAR but  
act: need to copy messages, otherwise could stuck in an endless loop

## 5.4 Haskell<sup>1</sup>

We initially thought that Haskell would be suitable best for just implementing the Par-Strategy after implementing all the strategies in it we found out that Haskell is extremely well suited to implement all the strategies.

We think this stems from the following facts: no side-effects (unless reflected in the types): is a must-have for STM, although it makes things more difficult in the beginning, in the end it turns out to be a blessing because one can guarantee that side-effects won't occur. We have taken care that the agents all run in side-effect free code.

STM: implementing concurrency is a piece-of-cake

extremely powerful static typesystem: in combination with side-effect free this results in the semantics of an update-strategy to be reflected in the Agent-Transformer function and the messaging-interface. This means a user of this approach can be guided by the types and can't abuse them. Thus the lesson learned here is that *if one tries to abuse the types of the agent-transformer or work around, then this is an indication that the update-strategy one has selected does not match the semantics of the model one wants to implement*. If this happens in Java, it is much more easier to work around by introducing global variables or side-effects but this is not possible in Haskell.

Thus our conclusion in using Haskell is that it is an extremely underestimated language in ABM/S which should be explored much more as we have shown that it really shines in this context and we believe that it could be pushed further even more.

but still it is not suitable for big models with heterogeneous agents, there things are lacking: see further research

note that the difference between SEQ and PAR in Haskell is in the end a 'fold' over the agents in the case of SEQ and a 'map' in the case of PAR dont have objects with methods which can call between each other but we need some way of representing agents. this is done using a struct type with a behaviour function and messaging mechanisms. important: agents are not carried around but messages are sent to a receiver identified by an id. This is also a big difference to java where don't really need this kind of abstraction due to the use of objects and their 'messaging'. messaging mechanisms have up- and downsides, elaborate on it.

concurrency and actors extremely elegant possible through: STM which only possible in languages without side-effect

the conc-version and the act-version of the agent-implementations look EX-ACTLY the same BUT we lost the ability to step the simulation!!!

This is the process of implementing the behaviour of the Agent as specified in the model. Although there are various kinds of Agent-Models like BDI but the basic principle is always the same: sense the environment, process messages, execute actions: change environment, send messages. According to [7] and also

---

<sup>1</sup>Code available under

<https://github.com/thalerjonathan/phd/tree/master/coding/papers/iteratingABM/haskell>

influenced by Actors from [2] one can abstract the abilities in each step of an Agent to be the following:

1. Process received messages
2. Create new Agents
3. Send messages to other Agents
4. Sense (read) the environment
5. Influence (write) the environment

## 5.5 Scala with Actors

direct support for actors

## 6 Conclusion

- Selecting the correct Iteration-Strategy is of most importance and must match the model semantics
- Java: parallelism and concurrency very easy due to elegant and powerful built-in synchronization primitives, high-performance and large number of agents possible due to aliasing and side-effects
- Surprise: Haskell can faithfully implement all strategies equally well, something not anticipated in the beginning. But: still much work to do to capture large and complex models (see further research), performance is a big issue but this has not been about performance (2000 Agents are enough)
- Scala with Actors: extremely elegant solutions possible

## 7 Further Research

### 7.1 Functional Reactive Programming

The implemented framework in Haskell is lacking features like TODO and is basically an attempt of reinventing Functional Reactive Programming (FRP). We were aware of the existence of this paradigm, especially the library Yampa, but decided to leave that one to a side and really keep our implementation clear and very basic. The next step would be to fusion our implementations with Yampa thus leveraging both approaches from which we hypothesize to gain the ability to develop much more complex models with heterogeneous agents.

## **7.2    Develop a small modelling-language which is close to the Haskell-Version of modelling agents therefore specification and implementation match**

See paper

## **7.3    Reasoning in Haskell about the Model & Simulation**

TODO: sketch ideas

## References

- [1] A, M. L., B, R. D., AND B, K. R. *A Framework for Megascale Agent Based Model Simulations on the GPU*. 2008.
- [2] AGHA, G. *Actors: A Model of Concurrent Computation in Distributed Systems*. MIT Press, Cambridge, MA, USA, 1986.
- [3] BOTTA, N., MANDEL, A., AND IONESCU, C. Time in discrete agent-based models of socio-economic systems. Documents de travail du Centre d'Economie de la Sorbonne 10076, Université Panthéon-Sorbonne (Paris 1), Centre d'Economie de la Sorbonne, 2010.
- [4] DAWSON, D., SIEBERS, P. O., AND VU, T. M. Opening pandora's box: Some insight into the inner workings of an Agent-Based Simulation environment. In *2014 Federated Conference on Computer Science and Information Systems* (Sept. 2014), pp. 1453–1460.
- [5] HUBERMAN, B. A., AND GLANCE, N. S. Evolutionary games and computer simulations. *Proceedings of the National Academy of Sciences* 90, 16 (Aug. 1993), 7716–7718.
- [6] NOWAK, M. A., AND MAY, R. M. Evolutionary games and spatial chaos. *Nature* 359, 6398 (Oct. 1992), 826–829.
- [7] WOOLDRIDGE, M. *An Introduction to MultiAgent Systems*, 2nd ed. Wiley Publishing, 2009.
- [8] YUXUAN, J. *The Agent-based Simulation Environment in Java*. PhD thesis, University Of Nottingham, School Of Computer Science, 2016.